

RISC-V RV32C: Compressed Instructions

Sparsh Mittal

Why add compressed instructions

The smaller the program, the smaller the area on a chip needed for the program memory, which can be a significant cost for embedded devices.

Indeed, that issue inspired ARM architects to retroactively add shorter instructions in the Thumb and Thumb-2 ISAs.

Smaller programs also lead to fewer misses in instruction caches, which saves power since off-chip DRAM accesses use much more energy than on-chip SRAM accesses, and improves performance.

RV32C

- **Prior ISAs** significantly expanded the number of instructions and instruction formats to shrink code size: adding short instructions with two operands instead of three, small immediate fields, and so on.
- **RV32C** takes a novel approach: every short instruction must map to one single standard 32-bit RISC-V instruction.
- Only the assembler and linker are aware of the 16- bit instructions, and it is up to them to replace a wide instruction with its narrow cousin.

- The compiler writer and assembly language programmer need not worry about RV32C instructions.
 - RV32I instructions are indistinguishable in RV32IC.
- A decoder translates all 16-bit instructions into their equivalent 32-bit version **before** they execute.
- This makes compression instructions inexpensive.

RV32C

Integer Computation

c.add { immediate }

c.add immediate * 16 to stack pointer

c.add immediate * 4 to stack pointer nondestructive

c.subtract

c. { shift left logical
shift right arithmetic
shift right logical } immediate

c.and { immediate }

c.or

c.move

c.exclusive or

c.load { upper } immediate

Loads and Stores

c. { float } { load
store } word { using stack pointer }

c.float { load
store } doubleword { using stack pointer }

Control transfer

c.branch { equal
not equal } to zero

c.jump { and link }




c.jump { and link } register

Other instructions

c.environment break

The immediate fields of the shift instructions and c.addi4spn are zero extended and sign extended for the other instructions.

Comparison of original and compressed instruction



Which instructions to compress?

- Three key observations
 1. Ten popular registers (a0–a5, s0–s1, sp, and ra) are accessed far more than the rest.
 2. Many instructions overwrite one of their source operands.
 3. Immediate operands tend to be small, and some instructions favor certain immediates.
 - E.g., loads and stores use only unsigned offsets in multiples of the operand size.

RVC Register Number

Integer Register Number

Integer Register ABI Name

Floating-Point Register Number

Floating-Point Register ABI Name

000	001	010	011	100	101	110	111
x8	x9	x10	x11	x12	x13	x14	x15
s0	s1	a0	a1	a2	a3	a4	a5
f8	f9	f10	f11	f12	f13	f14	f15
fs0	fs1	fa0	fa1	fa2	fa3	fa4	fa5

RV32C code for Insertion Sort and DAXPY

Instructions and code size for Insertion Sort and DAXPY for compressed ISAs

Benchmark	ISA	ARM Thumb-2	microMIPS	x86-32	RV32I+RVC
Insertion Sort	Instructions	18	24	20	19
	Bytes	46	56	45	52
DAXPY	Instructions	10	12	16	11
	Bytes	28	32	50	28

RV32C code for Insertion Sort

```
# RV32C (19 instructions, 52 bytes)
# a1 is n, a3 points to a[0], a4 is i, a5 is j, a6 is x
0: 00450693 addi a3,a0,4 # a3 is pointer to a[i]
4: 4705      c.li a4,1    # (expands to addi a4,x0,1) i = 1
Outer Loop:
6: 00b76363 bltu a4,a1,c # if i < n, jump to Continue Outer loop
a: 8082      c.ret                # (expands to jalr x0,ra,0) return from function
Continue Outer Loop:
c: 0006a803 lw a6,0(a3) # x = a[i]
10: 8636     c.mv a2,a3   # (expands to add a2,x0,a3) a2 is pointer to a[j]
12: 87ba     c.mv a5,a4   # (expands to add a5,x0,a4) j = i
InnerLoop:
14: ffc62883 lw a7,-4(a2) # a7 = a[j-1]
18: 01185763 ble a7,a6,26 # if a[j-1] <= a[i], jump to Exit InnerLoop
1c: 01162023 sw a7,0(a2)  # a[j] = a[j-1]
20: 17fd     c.addi a5,-1  # (expands to addi a5,a5,-1) j--
22: 1671     c.addi a2,-4  # (expands to addi a2,a2,-4)decr a2 to point to a[j]
24: fbe5     c.bnez a5,14  # (expands to bne a5,x0,14)if j!=0,jump to InnerLoop
Exit InnerLoop:
26: 078a     c.slli a5,0x2 # (expands to slli a5,a5,0x2) multiply a5 by 4
28: 97aa     c.add a5,a0   # (expands to add a5,a5,a0)a5 = byte address of a[j]
2a: 0107a023 sw a6,0(a5)  # a[j] = x
2e: 0705     c.addi a4,1   # (expands to addi a4,a4,1) i++
30: 0691     c.addi a3,4   # (expands to addi a3,a3,4) incr a3 to point to a[i]
32: bfd1     c.j 6         # (expands to jal x0,6) jump to Outer Loop
```

The twelve 16-bit instructions make the code 32% smaller.

RV32C instructions (starting with c.) are shown explicitly in this example, but normally assembly language programmers and compilers cannot see them.

RV32DC code for DAXPY

```
# RV32DC (11 instructions, 28 bytes)
# a0 is n, a1 is pointer to x[0], a2 is pointer to y[0], fa0 is a
0: cd09      c.beqz a0,1a      # (expands to beq a0,x0,1a) if n==0, jump to Exit
2: 050e      c.slli a0,a0,0x3   # (expands to slli a0,a0,0x3) a0 = n*8
4: 9532      c.add a0,a2        # (expands to add a0,a0,a2) a0 = address of x[n]
Loop:
6: 2218      c.fld fa4,0(a2)    # (expands to fld fa4,0(a2) ) fa5 = x[]
8: 219c      c.fld fa5,0(a1)    # (expands to fld fa5,0(a1) ) fa4 = y[]
a: 0621      c.addi a2,8        # (expands to addi a2,a2,8) a2++ (incr. ptr to y)
c: 05a1      c.addi a1,8        # (expands to addi a1,a1,8) a1++ (incr. ptr to x)
e: 72a7f7c3  fmadd.d fa5,fa5,fa0,fa4 # fa5 = a*x[i] + y[i]
12: fef63c27 fsd fa5,-8(a2)     # y[i] = a*x[i] + y[i]
16: fea618e3 bne a2,a0,6        # if i != n, jump to Loop
Exit:
1a: 8082     ret               # (expands to jalr x0,ra,0) return from function
```

The eight 16-bit instructions
shrink the code by 36%

Example 1 of compressed instruction

The assembler replaced the following 32-bit RV32I instruction:

```
addi a4,x0,1          # i = 1
```

with this 16-bit RV32C instruction:

```
c.li a4,1             # (expands to addi a4,x0,1) i = 1
```

The RV32C load immediate instruction is narrower because it must specify only one register and a small immediate.

Example 2 of compressed instruction

The assembler replaced the following 32-bit RV32I instruction:

`add a2,x0,a3` `# a2 is pointer to a[j]`

with this 16-bit RV32C instruction:

`c.mv a2,a3` `# (expands to add a2,x0,a3) a2 is pointer to a[j]`

The RV32C move instruction is merely 16 bits long because it specifies only two registers

Compressed 16-bit RVC instruction formats.

Format	Meaning	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0		
CR	Register	funct4				rd/rs1				rs2				op					
CI	Immediate	funct3		imm		rd/rs1				imm				op					
CSS	Stack-relative Store	funct3		imm						rs2				op					
CIW	Wide Immediate	funct3		imm										rd'		op			
CL	Load	funct3		imm			rs1'		imm		rd'		op						
CS	Store	funct3		imm			rs1'		imm		rs2'		op						
CB	Branch	funct3		offset			rs1'		offset				op						
CJ	Jump	funct3		jump target												op			

RV32C opcode map (bits[1 : 0] = 01)

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
000			nzimm[5]			0				nzimm[4:0]					01	CI c.nop
000			nzimm[5]			rs1/rd≠0				nzimm[4:0]					01	CI c.addi
001			imm[11 4 9:8 10 6 7 3:1 5]											01	CJ c.jal	
010			imm[5]			rd≠0				imm[4:0]					01	CI c.li
011			nzimm[9]			2				nzimm[4 6 8:7 5]					01	CI c.addi16s
011			nzimm[17]			rd≠{0, 2}				nzimm[16:12]					01	CI c.lui
100			nzuimm[5]		00	rs1'/rd'				nzuimm[4:0]					01	CI c.srli
100			nzuimm[5]		01	rs1'/rd'				nzuimm[4:0]					01	CI c.srai
100			imm[5]		10	rs1'/rd'				imm[4:0]					01	CI c.andi
100			0		11	rs1'/rd'			00	rs2'				01	CR c.sub	
100			0		11	rs1'/rd'			01	rs2'				01	CR c.xor	
100			0		11	rs1'/rd'			10	rs2'				01	CR c.or	
100			0		11	rs1'/rd'			11	rs2'				01	CR c.and	
101			imm[11 4 9:8 10 6 7 3:1 5]											01	CJ c.j	
110			imm[8 4:3]			rs1'				imm[7:6 2:1 5]					01	CB c.beqz
111			imm[8 4:3]			rs1'				imm[7:6 2:1 5]					01	CB c.bnez

rd', rs1', and rs2' refer to the 10 popular registers a0–a5, s0–s1, sp, and ra.

RV32C opcode map (bits[1 : 0] = 00)

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
000		0									0		00		CIW <i>Illegal instruction</i>	
000		nzuimm[5:4 9:6 2 3]									rd'		00		CIW c.addi4spn	
001		uimm[5:3]			rs1'			uimm[7:6]			rd'		00		CL c.fld	
010		uimm[5:3]			rs1'			uimm[2 6]			rd'		00		CL c.lw	
011		uimm[5:3]			rs1'			uimm[2 6]			rd'		00		CL c.flw	
101		uimm[5:3]			rs1'			uimm[7:6]			rs2'		00		CL c.fsd	
110		uimm[5:3]			rs1'			uimm[2 6]			rs2'		00		CL c.sw	
111		uimm[5:3]			rs1'			uimm[2 6]			rs2'		00		CL c.fsw	

RV32C opcode map (bits[1 : 0] = 10)

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
000			nzuimm[5]				rs1/rd \neq 0				nzuimm[4:0]				10	CI c.slli
000			0				rs1/rd \neq 0				0				10	CI c.slli64
001			uimm[5]				rd				uimm[4:3 8:6]				10	CSS c.fldsp
010			uimm[5]				rd \neq 0				uimm[4:2 7:6]				10	CSS c.lwsp
011			uimm[5]				rd				uimm[4:2 7:6]				10	CSS c.flwsp
100			0				rs1 \neq 0				0				10	CJ c.jr
100			0				rd \neq 0				rs2 \neq 0				10	CR c.mv
100			1				0				0				10	CI c.ebreak
100			1				rs1 \neq 0				0				10	CJ c.jalr
100			1				rs1/rd \neq 0				rs2 \neq 0				10	CR c.add
101			uimm[5:3 8:6]								rs2				10	CSS c.fsdsp
110			uimm[5:2 7:6]								rs2				10	CSS c.swsp
111			uimm[5:2 7:6]								rs2				10	CSS c.fswsp

References

<https://fprox.substack.com/p/riscv-c-extension>